

```
# Step 1: Import Libraries and Load Dataset
```

```
import pandas as pd
```

```
import numpy as np
```

```
from zipfile import ZipFile
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler, LabelEncoder
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
from sklearn.cluster import KMeans
```

```
from sklearn.metrics import silhouette_score
```

```
from sklearn.decomposition import PCA
```

```
import gym
```

```
from gym import spaces
```

```
import random
```

```
from collections import defaultdict
```

```
# Unzip and Load the dataset
```

```
file_path = 'diabetes+130-us+hospitals+for+years+1999-2008.zip'
```

```
with ZipFile(file_path, 'r') as zip_ref:
```

```
    csv_filename = zip_ref.namelist()[0]
```

```
    with zip_ref.open(csv_filename) as file:
```

```
        data = pd.read_csv(file)
```

```
# Display basic dataset information
```

```
print("Dataset Information:")
```

```
print(data.info())
```

```
print("\nFirst 5 Rows:")
```

```
print(data.head())
```

```

# Step 2: Data Preprocessing and Target Separation

# Define the target column
target_column = 'readmitted' # Change if needed

X = data.drop(columns=[target_column])
y = data[target_column]


# Encode categorical columns
X = pd.get_dummies(X, drop_first=True)


# Encode the target if it's categorical
if y.dtype == 'object':
    label_encoder = LabelEncoder()
    y = label_encoder.fit_transform(y)


# Split data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Standardize the data
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)


# Step 3: Classification Using Random Forest

# Train Random Forest Classifier
rf_classifier = RandomForestClassifier(random_state=42)
rf_classifier.fit(X_train, y_train)


# Predictions and metrics
y_pred = rf_classifier.predict(X_test)

print("\nRandom Forest Classifier Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))

```

```
# Confusion Matrix
```

```
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
plt.figure(figsize=(8, 6))
```

```
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
```

```
plt.title("Confusion Matrix for Random Forest Classifier")
```

```
plt.xlabel("Predicted")
```

```
plt.ylabel("Actual")
```

```
plt.show()
```

```
# Step 4: Clustering Using K-Means
```

```
# K-Means Clustering
```

```
kmeans = KMeans(n_clusters=3, random_state=42)
```

```
kmeans.fit(X_train)
```

```
clusters = kmeans.predict(X_train)
```

```
silhouette_avg = silhouette_score(X_train, clusters)
```

```
print("\nSilhouette Score for K-Means Clustering:", silhouette_avg)
```

```
# PCA for Visualization
```

```
pca = PCA(n_components=2)
```

```
X_pca = pca.fit_transform(X_train)
```

```
plt.figure(figsize=(10, 8))
```

```
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=clusters, cmap='viridis', s=50, alpha=0.6)
```

```
plt.title("K-Means Clustering Visualization (PCA-reduced data)")
```

```
plt.xlabel("PCA Component 1")
```

```
plt.ylabel("PCA Component 2")
```

```
plt.colorbar(label="Cluster")
```

```
plt.show()
```

```
# Step 5: Reinforcement Learning Environment Setup and Q-Learning
```

```
# Define a custom gym environment for feature selection
```

```

class DiabetesEnv(gym.Env):
    def __init__(self, data, target):
        super(DiabetesEnv, self).__init__()
        self.data = data.select_dtypes(include=[np.number]) # Only numeric features
        self.target = target
        self.action_space = spaces.Discrete(self.data.shape[1]) # One action per feature
        self.observation_space = spaces.Box(low=0, high=1, shape=(self.data.shape[1],))
        self.reset()

    def reset(self):
        self.state = np.random.choice(self.data.index)
        return self.data.iloc[self.state].values

    def step(self, action):
        reward = np.corrcoef(self.data.iloc[:, action], self.target)[0, 1]
        done = True # Single-step environment
        return self.data.iloc[self.state].values, reward, done, {}

# Initialize the environment
env = DiabetesEnv(pd.DataFrame(X_train), y_train)

# Q-Learning parameters
num_episodes = 100
learning_rate = 0.1
discount_factor = 0.9
epsilon = 0.1 # Epsilon-greedy

# Q-table
q_table = defaultdict(lambda: np.zeros(env.action_space.n))

for episode in range(num_episodes):

```

```

state = env.reset()

done = False

while not done:

    if random.uniform(0, 1) < epsilon:

        action = env.action_space.sample()

    else:

        action = np.argmax(q_table[state.tobytes()])

    next_state, reward, done, _ = env.step(action)

    best_next_action = np.argmax(q_table[next_state.tobytes()])

    td_target = reward + discount_factor * q_table[next_state.tobytes()][best_next_action]

    td_error = td_target - q_table[state.tobytes()][action]

    q_table[state.tobytes()][action] += learning_rate * td_error

    state = next_state


# Display sample Q-values after training
sample_q_values = dict(list(q_table.items())[:5])

print("\nSample Q-table values after training:", sample_q_values)


# Graph of rewards per action
actions = list(range(env.action_space.n))

rewards = [np.corrcoef(env.data.iloc[:, a], env.target)[0, 1] for a in actions]

plt.figure(figsize=(10, 6))

plt.bar(actions, rewards, color='blue', alpha=0.7)

plt.xlabel("Action (Feature Index)")

plt.ylabel("Correlation with Target (Reward)")

plt.title("Reward per Action (Feature)")

plt.show()

```